

# Case Study: LLM-Driven SQL Optimization for Upwork

## Project Summary

At Upwork, I led a large-scale initiative to optimize over **4,000 legacy SQL queries** using **LLMs integrated with PROM (Prompt-Oriented Modeling)**. This project resulted in **six-figure monthly savings** by identifying inefficiencies in query design, reducing data volume scanned, and minimizing Snowflake credit usage.

---

## Problem

Upwork's data engineering team was managing thousands of legacy SQL files across multiple marketing and analytics pipelines. These queries were: - Poorly optimized for Snowflake, leading to high compute costs - Repetitive or redundant - Not scalable as query volumes grew - Lacking visibility into cost drivers or optimization opportunities

The result was a **significant increase in infrastructure spending** and difficulty scaling the data platform efficiently.

---

## Goal

- Analyze 4,000+ SQL files
  - Evaluate their cost, data volume scanned, and compute credit usage
  - Use an **LLM (via PROM)** to generate optimized versions of each query
  - Quantify and deliver potential savings
  - Provide recommendations to engineers and analysts on improvement
- 

## Solution & Technical Approach

- **Query Ingestion & Parsing:** Built a system to iterate over 4,000+ SQL files and ingest each query for analysis.
  - **Cost Modeling:** For each query, the system retrieved metadata including:
    - Execution time
    - Data scanned (TB/GB)
    - Snowflake credit consumption
  - **LLM Optimization Layer:** Integrated PROM to interpret each query's logic and generate **optimized SQL alternatives** while preserving functional intent.
  - **Cost Simulation:** Estimated savings by comparing pre- and post-optimization metrics.
  - **Recommendation Engine:** Outputted a report per file, including:
    - Optimization suggestions
    - Line-by-line diffs
    - Estimated monthly savings
    - Engineering action level (low, medium, high priority)
-

## Results

- **Analyzed:** 4,000+ SQL queries
  - **Optimized:** Over 80% of files had actionable improvements
  - **Savings:** Project enabled **six-figure monthly cost reductions**
  - **Efficiency:** Cut down engineering review time by 70% with automation
  - **Scalability:** The system is now extendable to future pipelines and integrated into query review processes
- 

## Impact

- Empowered the analytics and data engineering teams to **reduce query cost and complexity at scale**
  - Allowed leadership to **monitor cost-saving metrics** per team or pipeline
  - Established a framework for **future LLM-based DevOps tooling** at Upwork
- 

## Tools & Technologies

- **PROM (LLM framework)**
  - **Snowflake, SQL, Python**
  - **Airflow** for orchestration
  - **GitHub Actions** for automation
  - **Looker** for reporting and stakeholder dashboards
- 

## Lessons Learned

- LLMs can significantly improve technical workflows when paired with **domain-specific metadata and constraints**
- Engineering teams benefit from cost visibility that is **proactive, not reactive**
- Optimization is not just about compute savings—it improves **performance, reliability, and developer experience**